

# Recitation 13: Random forests and artificial neural networks

Seung-hun Lee

Columbia University  
Introduction to Econometrics II Recitation

May 2nd, 2022

Random forest

# Classical random forests: Classification and regression trees

- We have an IID data  $(y_i, x_i)$  for  $i \in \{1, \dots, n\}$
- Decision trees represent classification and regression problems using nonlinear basis functions (indicators of categories or cutoff)
  - ▶ Branch: Path from one node to other
  - ▶ Split: Process of dividing a sample based on a criterion specified at a node; creates two branches per non-terminal nodes
  - ▶ Depth: Number of stages with splits
  - ▶ Leaves: Output; or nodes on the last stage
  - ▶ Pruning: Reducing the size of the tree (steps or fewer splits)
- The goal is to find an algorithm that best predicts out-sample properties

# CART uses recursive greedy algorithm

- Recursive in that there are repeated nodes with partitioning of two-way splits
  - ▶ Select a variable  $x_k$  and threshold/class that allows each split outcome to be as different as possible (between variance and information gain)
  - ▶ Between variance: Variation between each group mean to overall mean,  $R^2$  of

$$y = \beta \mathbb{I}(x_k > t) + u$$

- ▶ Information gain: For a potential split  $S = S_1 \cup S_2$

$$IG = L(S) - \left[ \frac{|S_1|}{|S|} L(S_1) + \frac{|S_2|}{|S|} L(S_2) \right], \quad L(C) = - \sum_{j=1}^J f_j \log f_j$$

i.e. splits  $S_1$  and  $S_2$  that has lower combined entropies (leading to high IG)

- Greedy in that at each non-terminal node, the best test selected is only locally optimal
  - ▶ No look-forward
  - ▶ Due to NP-hardness of the optimal tree problem, this is cheaper

# Finding the best CART is not easy

- Selection of hyperparameters (parameters whose value is used to control the learning process such as split criterion, stopping criteria)
  - ▶ Can lead to overfitting and bad out-sample predictions if improperly selected
- Pruning: Making the trees more compact by eliminating some splits
  - ▶ Cross-validation is used
- Greedy algorithm: Decision tree may not be stable and the most optimal possible, especially when the top node is incorrectly selected

# Bootstrap AGGREGATING, or Bagging

- Like ensemble method, it averages from multiple decision trees to enhance stability and accuracy of prediction, classification, and regression
- By averaging, bagging protects from overfitting the data and prevents instability
- The procedure is very similar to that of bootstrapping. They are as follows
  1. Bootstrap: Let's have a sufficiently large  $B$  resamples with replacement. In each resampling step, we fit a tree.
  2. Aggregation: We average the predictions of the  $B$  trees to predict  $y$  given  $x$ .
- Predictive capability: Are the predictions less dispersed? Then this is good prediction

# Random Forest

- Averaging multiple deep decision trees, trained on different parts of the same training data set
- Unlike bagging, random forests choose splits from a smaller number of randomly selected covariates at each step of the tree construction (feature bagging)
  - ▶ With  $k$  available features,  $\sqrt{k}$  is used in classification and  $\max\{k/3, 5\}$  in regression
- Less correlation among trees compared to earlier methods (variety of top splits)
- Less overfitting and more stability
  - ▶ Variable with huge decline in split-criterion (out-of-bag-error from permutation over  $B$  trees) can be considered important
- Key problem is loss of interpretability (black-box procedure in selection of covariates)

## Causal forest (Wagner, Athey 2018)

- Random forest to estimate treatment effect (under CIA)
- We first estimate the propensity score  $\hat{p}_i$  and transform  $y_i$  into

$$\hat{y}_i = \frac{D_i - \hat{p}_i}{\hat{p}_i(1 - \hat{p}_i)} y_i$$

which is an IPW formula

- By taking the expected value of  $\hat{y}_i$  we get

$$\widehat{TE}(x) = E[\hat{y}_i|x]$$

- We then do 'honest' inference (defined in the paper) by estimating on a subsample that is different from the training and test subsamples (double-sample trees)
- We can also do normal inference



## Local linear forest (Friedberg, Tibshirani, Athey, Wagner 2020)

- Random forests are used to generate weights that can serve as a kernel for a local linear regression and estimates for the conditional mean function
- Think of random forests as adaptive weight generators

$$\hat{m}(x_0) = \frac{1}{B} \sum_{i=1}^n \sum_{b=1}^B y_i \frac{\mathbb{I}[X_i \in L_b(x_0)]}{|L_b(x_0)|} = \sum_{i=1}^n y_i \underbrace{\frac{1}{B} \sum_{b=1}^B \frac{\mathbb{I}[X_i \in L_b(x_0)]}{|L_b(x_0)|}}_{\alpha_i(x_0)}$$

- $L_b(x_0)$  represents the leaf at  $b$ 'th tree,  $|\cdot|$  indicates the size of the leaves.
- $\alpha_i(x_0)$  is the fraction of trees in which an observation appears in the same leaf as the target value of the covariate vector

## Local linear forest (Friedberg, Tibshirani, Athey, Wagner 2020)

- Local linear forests solve the weighted least problem below with penalty term using the weights defined above

$$\begin{pmatrix} \hat{m}(x_0) \\ \hat{\theta}(x_0) \end{pmatrix} = \arg \min_{m, \theta} \left[ \sum_{i=1}^n \alpha_i(x_0) (y_i - m(x_0) - (x - x_0)\theta(x_0))^2 + \lambda \|\theta(x_0)\|^2 \right]$$

- Performs better when smooth, strong signals are present (random forests don't do so well in this case)
- Also overcomes curse of dimensionality

# Gradient-boosted trees

- Improve the performance of weak learners by boosting the weight of observations they misspecify
- Start with a loss function  $\min \sum_{i=1}^n L(y_i, \hat{y}_i)$ , say,  $(y_i - \hat{y}_i)^2$
- $v(x, S)$ : Average value of  $y$  in the leaf that contains  $x$  in a shallow tree with splits  $S$
- Then we proceed in these steps
  1. Begin with an initial prediction:  $\hat{y}_i^{(0)} = E[y_i]$
  2. For each iteration  $b \in \{1, \dots, B\}$ , we choose split  $S_b$  in the tree of depth  $d$  to minimize

$$\sum_{i=1}^n L(y_i, \hat{y}_i^{(b-1)} + v(x_i, S))$$

and update our prediction  $\hat{y}_i^{(b)} = \hat{y}_i^{(b-1)} + \epsilon v(x_i, S_b)$ . Note that  $\epsilon$  is the learning rate and is not  $\epsilon = 1$ , which would mean updating the prediction by the exact value and risk overshooting

## Other ways to do gradient-boosted trees

- We can also tune the model by including a complexity penalty in the loss function. So we minimize

$$\sum_{i=1}^n L(y_i, \hat{y}_i^{(b-1)} + v(x_i, S)) + C(S)$$

- XGBoost: Extreme gradient-boosting, increases the speed of the procedure using quadratic loss and complexity penalty functions


$$\sum_{i=1}^n (G_i v(x_i, S) + \frac{1}{2} H_i v(x_i, S)^2) + \gamma K + \frac{\lambda}{2} \sum_{k=1}^K v_k^2$$

where  $G_i$  and  $H_i$  are each the first and second derivatives of the  $L(y_i, \hat{y}_i)$  w.r.t  $\hat{y}_i$  at  $(y_i, \hat{y}_i^{(b-1)})$ ,  $K$  is the number of leaves, and  $v_k$  is the predictor on leaf  $k$

- Tuning and some hyperparameters still need to be determined

# Artificial neural networks

# Architecture of artificial neural networks



# Architecture of artificial neural networks

- Input layers have covariates  $x_m$
- Lines between layers have weights  $w$
- Each direction represents signal  $\sigma(\cdot)$ , known as activation function
  - ▶ A sigmoid (logit)
  - ▶ Multinomial logit (softmax)
  - ▶ Rectified linear unit (ReLU,  $\sigma(t) = \max\{0, t\}$ )

## Single layer perceptions: Setup and forward pass

- $M$  neurons in input,  $L$  nodes in hidden layer
- Input neurons are denoted  $x_1, \dots, x_M$  have weights  $w_{1lm}$ ,  $m \in \{1, \dots, M\}$ ,  $l \in \{1, \dots, L\}$ , and emit  $\sigma(\sum_m w_{1lm}x_m)$  to a hidden layer  $l$
- Hidden layer has weight for each node,  $w_{2l}$ , used to combine inputs from previous nodes to generate a prediction for the output node  $\hat{y} = g(\sum_l w_{2l}\sigma(\sum_m w_{1lm}x_m))$
- Epoch: One forward pass + backward pass
- Forward pass: Given loss function  $L(y, \hat{y}(\theta))$  and an initial value  $\theta = (w_1, w_2)$ , compute

$$z_{il} = \sigma(w'_{1l}x_i)$$

and then the output can be produced as

$$\hat{y}_i(\theta) = w'_2 z_i = \sum_l w_{2l} \sigma \left( \sum_m w_{1lm} x_{im} \right)$$



## Single layer perceptions: Backward pass

- Compute the components of the loss functions, then take gradients with respect to elements in  $\theta$ , and do approximate Newton iteration (e.g. Minibatch gradient descent)

$$\theta^{(s+1)} = \theta^{(s)} - \epsilon_s \frac{\partial L}{\partial \theta}(\theta^{(s)})$$

- $\epsilon_s$  is the learning rate which ideally should goes to 0
- Backpropagation: An automated differentiation and application of chain rule to identify gradients used to update the values of the weights

# Deep learning: Multiple hidden layers

- Many (hyper)parameters, activation functions to choose
- We have  $p$  covariates,  $K$  outputs,  $D$  hidden layers, and  $M$  nodes per layer ( $p, K$  given)
- In total, we have  $pM + (D - 1)M^2 + kM$  parameters (or weights) to work with
- Optimal  $D$  and  $M$ ? It is known that deep artificial neural networks are preferable to wider ones
  - ▶ Small  $D$  does not fit well, large  $D$  induces less overfitting problem especially if there are many small weights (also, we gain on expressiveness)
  - ▶ Incorporating penalty term  $\lambda ||\theta||_q^q$  helps too
- Dropout: Randomly eliminates some links and thins the neural networks, which is then retrained and tested/compared with all the eliminated units brought back in
  - ▶ Realistic alternative to cross-validation

# Double machine learning debiasing techniques

- ATE under CIA

$$y_i = \mu(X_i, D_i) + \epsilon_i(D_i), \quad D_i = \mathbb{I}(u_i < p(X_i)), \quad (\epsilon(0), \epsilon(1)) \perp\!\!\!\perp u|X$$

- Use ML to estimate the  $p(X_i)$  and  $\mu(X_i, D_i)$  and obtain IPW ATE
- Debiasing?  $p(X_i)$  and  $\mu$  are nuisance parameters whose errors affect TE estimates
- Decouple the TE from the  $p(X_i)$  and  $\mu$  using Neyman orthogonalization (block-diagonalize the (Fisher) information matrix) - ML method involved
  - ▶ Split data to  $K$  folds, and leave one out
  - ▶ Estimate  $\hat{p}_k = \Pr(D = 1|X)$  using LASSO or other ML methods
  - ▶ Take the Neyman-orthogonalized combination of the estimates of  $p(X_i)$  and  $\mu$
  - ▶ If functions are smooth, estimates have standard asymptotics and we can do standard tests afterwards